

# Numerical Methods in Julia

## Course 2 — Short Macroeconomics Course Using Julia

November 26, 2025

# Outline

- 1 Overview
- 2 Performance and Benchmarking
- 3 Numbers and Errors
- 4 Linear Algebra Essentials
- 5 Root Finding
- 6 Optimization
- 7 Differentiation and Integration
- 8 Interpolation
- 9 Ordinary Differential Equations
- 10 Randomness and Monte Carlo
- 11 Wrap-Up

- Workhorse numerical tools for macro: linear algebra, root finding, optimization, integration, interpolation, ODEs, and Monte Carlo.
- Julia idioms: broadcasting, in-place updates, types, and performance.
- Reproducibility: environments, RNG, and benchmarking.

# Guiding Principles

- Stability and conditioning.
- Prefer library functions (e.g. factorizations) to manually built functions.
- Execution time: use `BenchmarkTools.jl`.

# Matrix operations: loops vs. broadcast (elementwise)

**Goal:** add a scaled matrix into another,  $C \leftarrow a \cdot A + C$  (AXPY on matrices).

```
using LinearAlgebra
# Broadcasting (concise)
C .+= a .* A
# Full broadcast (no temporaries): equivalent but preferred
@. C = a*A + C           # or: C .+= a .* A
# Loop
function axpy_loop!(C, a, A)
    @inbounds @simd for j in axes(C,2), i in axes(C,1)
        C[i,j] += a * A[i,j]
    end
    return C
end
```

**Notes.** Broadcasting with dotted operators fuses elementwise ops and avoids temporaries. Loops give maximal control (@inbounds to skip bounds checks, @simd to vectorize).

# BLAS and factorizations (dense linear algebra)

**Goal A:** matrix–matrix and matrix–vector products (use BLAS).

```
# High-level BLAS-backed ops
Y = A * x                      # gemv (matrix–vector)
C = A * B                      # gemm (matrix–matrix)
# In-place BLAS (avoid allocations)
mul!(Y, A, x)                  # Y := A*x
mul!(C, A, B, 1.0, 0.0)        # C := 1*A*B + 0*C
```

**Goal B:** solve many linear systems with the same  $A$  *without* re-factorizing.

```
# One-off solve (preferred over inv(A)*B)
X = A \
# Many solves with same A: factor once, reuse
F = lu(A)                      # or cholesky(A) if SPD
X = F \ B                       # multiple RHS columns
```

**Notes.** Never form  $\text{inv}(A)$  for solving; use  $\backslash$  or a reusable factorization ( $\text{lu}$ ,  $\text{cholesky}$ ,  $\text{qr}$ ) for speed and numerical stability.

# Broadcasting and In-place

```
x = rand(1_000_000)
y = similar(x)
y .= @. 2x + 1 # fused broadcast
```

# Avoid Allocations

```
function scale!(y, a, x)
    @inbounds @simd for i in eachindex(x)
        y[i] = a*x[i]
    end
    return y
end
```

## Why the extras help

- `scale!` — in-place update; no new array is created.
- `eachindex(x,y)` — fast, allocation-free indices that match shapes.
- `@inbounds` — skips bounds checks (you guarantee indices are valid).
- `@simd` — hints LLVM to auto-vectorize the loop for speed.

```
using BenchmarkTools
f(x) = sum(@. 2x + 1)
x = rand(10_000)
@btime f($x)
```

## Why the extras help

- Why the \$ in `@btime f($x)`:
- The dollar *uses* `x` into the benchmark.
- That means BenchmarkTools substitutes the current value of `x` directly into the tested expression so the benchmark doesn't include (1) global-variable lookup and (2) spurious allocations from treating `x` as a non-constant global.

# Type Stability

```
# inconsistent return types hurt performance
bad(x) = x > 0 ? 1 : 1.0
# fix by consistent types
good(x) = x > 0 ? 1.0 : 1.0
```

# Floating Point Essentials

```
using Printf
x = 0.1 + 0.2
@printf("%.17f\n", x) # 0.3000000000000004
isapprox(x, 0.3; atol=1e-12)
```

- Use `isapprox` with tolerances.

# Conditioning and Stability

```
A = [1.0 1.0; 1.0 1.000001]
condA = cond(A)    # 2-norm condition number
```

- Bad conditioning magnifies errors.
- Prefer stable algorithms (QR, SVD) over naive ones (which work, but are numerically unstable).

# BLAS-backed Operations

```
using LinearAlgebra
A = randn(3,3); b = randn(3)
A*b          # matrix-vector
A*A'         # matrix-matrix
```

# Solve, Do Not Invert

```
A = randn(3,3); b = randn(3)
x = A \ b          # solve Ax = b
# If repeated solves with same A, factorize once:
F = lu(A); x2 = F \ b
```

# Factorizations

```
F = qr(randn(4,3))
Q = Matrix(F.Q); R = F.R
S = svd(randn(4,3))
U, s, Vt = S.U, S.S, S.Vt
```

# Eigenvalues and Symmetry

```
A = Symmetric(randn(4,4)); A = A*A' # SPD  
D, V = eigen(A) # eigen decomposition
```

- `symmetric` is a function from `LinearAlgebra` telling Julia to treat an object as symmetric.
- enables safer and faster algorithms.

# Sparse Matrices

```
using SparseArrays
S = spdiagm(0 => fill(2.0,5), 1 => fill(-1.0,4), -1 => fill(-1.0,4))
b = ones(5)
x = S \ b
```

- sparse matrices: arrays where most entries are zero.
- A sparse matrix in Julia is usually CSC (Compressed Sparse Column) format.

# Fixed Point Iteration

```
# Solve x = cos(x)
x = 0.5
for k in 1:30
    x = cos(x)
end
x
```

- $f(x) = 0$  is mathematically equivalent to looking for a fixed point for which  $x = g(x)$  when we rewrite the problem as:
- $x = x - f(x)$ .
- To apply the function iteration, we need a starting point which must be reasonably close to the root of  $f(x)$ .

# Newton Method (Scalar)

```
# Solve f(x) = x^2 - 2 = 0
f(x) = x^2 - 2
fp(x) = 2x
x = 1.0
for k in 1:10
    x -= f(x)/fp(x)
end
```

- The algorithm of the Newton's method is pretty straightforward as the linearization is done using the Taylor expansion.
- Assume a function  $f$  and an initial guess of the root given by  $x_0$ . We create an iteration equation on  $x_{k+1}$  given  $x + k$  using a Taylor approximation of order one to replace  $f(x)$ :
- $f(x) \approx f(x_k) + f'(x_k)(x - x_k) = 0$  which gives the update formula:
- $x_{k+1} = x_k - \frac{f(x_k)}{f'(x_k)}$ .

# Newton Method (Vector)

```
using LinearAlgebra
function newton(F, J, x; iters=20)
    for k in 1:iters
        x -= J(x) \ F(x)
    end
    return x
end
```

# Unconstrained Optimization

```
# Gradient descent demo
f(x) = (x-2)^2
fp(x) = 2(x-2)
x = 0.0
for k in 1:20
    x -= 0.2*fp(x)
end
x
```

# Line Search Template

```
function linesearch(f, x; d=-1.0, alpha=1.0, beta=0.5)
    fx = f(x)
    while f(x + alpha*d) > fx
        alpha *= beta
    end
    return alpha
end
```

- A simple backtracking line search: start with step  $\alpha$
- repeatedly shrink it by  $\beta$  until moving from  $x$  to  $x + \alpha d$  actually lowers  $f$ , then return that step size.

# Quasi-Newton Sketch

```
# Illustrative only
function steepest_descent(f, g, x; steps=30)
    for k in 1:steps
        d = -g(x)
        a = 0.1
        x += a*d
    end
    return x
end
```

- Moves iteratively in the negative gradient direction  $d = -g(x)$  — the direction of steepest local decrease
- $a = 0.1$  each time, for steps iterations, to minimize  $f$

# Finite Differences

```
# Forward difference
fd(f, x, h=1e-6) = (f(x+h) - f(x)) / h
fd(sin, 0.3)
```

# Central Differences

```
cd(f, x, h=1e-6) = (f(x+h) - f(x-h)) / (2h)  
cd(exp, 0.0)
```

# Composite Simpson Rule

```
function simpson(f, a, b, n)
    n % 2 == 0 || error("n must be even")
    h = (b-a)/n
    s = f(a) + f(b)
    for k in 1:2:n-1
        s += 4*f(a + k*h)
    end
    for k in 2:2:n-2
        s += 2*f(a + k*h)
    end
    return s*h/3
end
simpson(x->exp(-x^2), 0.0, 1.0, 200)
```

# Piecewise Linear Interpolation

```
x = 0:0.5:5
y = @. sin(x) + 0.1*x
function lininterp(xg, yg, x)
    i = searchsortedlast(xg, x)
    i == length(xg) && (i -= 1)
    t = (x - xg[i])/(xg[i+1]-xg[i])
    return (1-t)*yg[i] + t*yg[i+1]
end
lininterp(collect(x), collect(y), 2.2)
```

- piecewise-linear interpolator
- given a sorted grid  $xg$  with values  $yg$ , it returns an approximate  $y$  at any  $x$  by joining neighboring points with straight lines.

# Using Package Interpolations

```
#use Interpolations
using Interpolations

#create some data
A_x = 1.0:4.0:50.0

A    = collect(2 .* A_x)

itp  = interpolate(A, BSpline(Cubic(Line()), OnGrid()))
sitp = scale(itp, A_x)

sitp(3.0)
```

- implements a variety of interpolation algorithms.
- supports B-splines and irregular grids.

# Explicit Euler

```
# x' = -x, x(0)=1
f(t,x) = -x
function euler(f, t0, x0, h, n)
    t = t0; x = x0
    for k in 1:n
        x += h*f(t, x)
        t += h
    end
    return x
end
euler(f, 0.0, 1.0, 0.01, 100)
```

# Explicit Euler

- **Forward Euler for**  $x'(t) = f(t, x)$ : starting at  $(t_0, x_0)$ , iterate

$$x_{k+1} = x_k + h f(t_k, x_k), \quad t_{k+1} = t_k + h,$$

for  $k = 0, \dots, n - 1$ .

- For  $f(t, x) = -x$ ,  $x(0) = 1$ ,  $h = 0.01$ ,  $n = 100$ , the method approximates

$$x(1) \approx (1 - 0.01)^{100} \approx 0.366,$$

close to the exact value  $e^{-1} \approx 0.368$ .

# Runge–Kutta 4

```
function rk4(f, t0, x0, h, n)
    t = t0; x = x0
    for k in 1:n
        k1 = f(t, x)
        k2 = f(t + h/2, x + h*k1/2)
        k3 = f(t + h/2, x + h*k2/2)
        k4 = f(t + h, x + h*k3)
        x += h*(k1 + 2k2 + 2k3 + k4)/6
        t += h
    end
    return x
end
rk4(f, 0.0, 1.0, 0.01, 100)
```

# Runge–Kutta 4 Algorithm

- **RK4 update for**  $x'(t) = f(t, x)$ : from  $(t_n, x_n)$  with step  $h$ ,

$$k_1 = f(t_n, x_n),$$

$$k_2 = f\left(t_n + \frac{h}{2}, x_n + \frac{h}{2}k_1\right),$$

$$k_3 = f\left(t_n + \frac{h}{2}, x_n + \frac{h}{2}k_2\right),$$

$$k_4 = f(t_n + h, x_n + h k_3),$$

$$x_{n+1} = x_n + \frac{h}{6}(k_1 + 2k_2 + 2k_3 + k_4), \quad t_{n+1} = t_n + h.$$

It samples the slope at four locations within the step and takes a weighted average, greatly improving accuracy over forward Euler.

- **Accuracy and use:** RK4 has a good accuracy–cost tradeoff for smooth  $f$ . It is explicit, easy to implement, and typically stable for moderate  $h$ .

# RNG Control

```
using Random  
Random.seed!(2025)  
rand(), randn()
```

# Monte Carlo Integration

```
# Estimate E[g(Z)] for Z ~ N(0,1)
using Statistics
function mc(g, N=10_000)
    z = randn(N)
    return mean(g.(z))
end
mc(z->z^2)
```

- **Monte Carlo integration:** The function `mc(g,N)` draws  $N$  iid samples  $Z_1, \dots, Z_N \sim \mathcal{N}(0, 1)$  and returns the sample mean  $\frac{1}{N} \sum_{i=1}^N g(Z_i)$ . By the Law of Large Numbers, this converges to  $\mathbb{E}[g(Z)]$  as  $N \rightarrow \infty$ . Broadcasting `g.(z)` applies  $g$  elementwise to the vector of draws.
- **Example:** `mc(z->z^2)` estimates  $\mathbb{E}[Z^2]$  for  $Z \sim \mathcal{N}(0, 1)$ , which equals 1 (the variance). With large  $N$ , the return should be close to 1; the error decays on the order of  $1/\sqrt{N}$ .

# Variance Reduction: Antithetic

```
function mc_anti(g, N=10_000)
z = randn(div(N,2))
return mean((g.(z) .+ g.(-z))/2)
end
```

- **Construction.** For  $Z \sim \mathcal{N}(0, 1)$ , draw  $z_i$  and pair it with its mirror  $-z_i$ ; estimate  $\mathbb{E}[g(Z)]$  by  $\hat{\mu} = \frac{1}{N} \sum_{i=1}^{N/2} [g(z_i) + g(-z_i)]$  (unbiased since  $Z \stackrel{d}{=} -Z$ ).
- **Why it helps.** The pair induces negative correlation when  $g$  is roughly monotone, giving  $\text{Var}\left(\frac{g(Z)+g(-Z)}{2}\right) = \frac{\text{Var}(g(Z))}{2}(1 + \rho)$  with  $\rho = \text{corr}(g(Z), g(-Z)) < 0$ .

# Summary

- Covered numerical methods.
- Emphasized stability, conditioning, and reproducibility.
- Next chapter: solving and simulating DSGE models.